

Homework 1: Getting Started with Probability

Instructions: Submit a single Jupyter notebook (.ipynb) of your work to Collab by 11:59pm on the due date. All code should be written in Python. **Be sure to show all the work involved in deriving your answers! If you just give a final answer without explanation, you may not receive credit for that question.**

You may discuss the concepts with your classmates, but write up the answers entirely on your own. Do not look at another student's answers, do not use answers from the internet or other sources, and do not show your answers to anyone. **Cite any sources you used outside of the class material (webpages, etc.), and list any fellow students with whom you discussed the homework concepts.**

1. Say 12% of the population is rich, 8% of the population is famous, and 6% of the population is both rich and famous. Define events R = "person is rich" and F = "person is famous" for some randomly selected person in the population. Write an expression for each of the following events using set operations involving the events R and F . **Here you can just give the answer, and do not need to show any work.**

Example: "The person is rich and famous" would be the event $R \cap F$.

- (a) The person is not famous.
 - (b) The person is rich but not famous.
 - (c) The person is either rich or famous (or both).
 - (d) The person is neither rich nor famous.
2. Calculate the probabilities for each of the four events in the previous problem. **Be sure to show your intermediate steps and list any probability rules that you use.**
Example: The probability for "The person is rich and famous" would be $P(R \cap F) = 0.06$.
 3. You have two black socks, two white socks, two red socks, and two blue socks in your sock drawer. If you pick two socks out at random (without looking!), what is the probability that they match? **Explain how you arrived at the answer, don't just give a number!**
 4. According to the American Lung Association, there is a 0.13% chance to develop lung cancer. Of the people who have lung cancer, 90% of them are smokers. In the population of people who do not have lung cancer, 16.9% are smokers.
 - (a) What percentage of the total population are smokers?
 - (b) If you are a smoker, what is your probability to develop lung cancer?
 - (c) If you are not a smoker, what is your probability to develop lung cancer?

5. In this exercise we will be using data from the OASIS brain database, a publicly-available resource:

<http://www.oasis-brains.org>

You will be classifying dementia from the volume of the hippocampus, a brain structure that is critical to memory. The data you will use is in the spreadsheet `OASIS-hippocampus.csv`, which you can download from the class website. The data consists of the hippocampal volume, derived from MRI, for elderly subjects, including healthy control subjects and those with mild to moderate dementia. Model the right hippocampal volume (`RightHippoVol`) as a normal random variable X_1 and the left hippocampal volume (`LeftHippoVol`) as a normal random variable X_2 . Then model the diagnosis (`Dementia`) as a binary random variable Y ($Y = 0$: healthy control, $Y = 1$: dementia). Use the training subset of the data (`TrainData = 1`) to learn the mean and variance parameters for a naïve Bayes classifier. Finally, apply your naïve Bayes classifier to get a probabilistic diagnosis of the test subset of the data (`TrainData = 0`).

Do the following:

- Plot the data as a 2D scatterplot (right and left hippocampal volume as the two axes). Use two different colors for the two classes (healthy/dementia). Do you think there is separation between the two classes?
- Plot two density plots for the left and right hippocampus volumes. Again, plot a different density for the two classes (with the same colors as your scatterplot).
- Run your classifier on the testing data. For each data point, if your classifier probability is greater than 0.5, predict that it is a dementia patient. Then compare with the actual label to see if your classifier is correct. Report your classifier accuracy on the testing data (number of correct classifications divided by size of the test data).

Note: There are Python machine learning libraries that include a method for naïve Bayes. You can't use these to solve this problem! You have to write your own code to implement naïve Bayes. However, it's okay to try out a library's version to check if it gives a similar answer to your code.